

Module 2, Part 1: Statistical Learning

TMA4268 Statistical Learning V2020

Stefanie Muff, Department of Mathematical Sciences, NTNU

January 10, 2020

Acknowledgements

- A lot of this material stems from Mette Langaas and her TAs (especiall Julia Debik). I would like to thank Mette for the permission to use her material!
- Some of the figures and slides in this presentation are taken (or are inspired) from James et al. (2013).

Introduction

Learning material for this module

- James et al (2013): An Introduction to Statistical Learning. Chapter 2.
- Additional material (in this module page) on random variables, covariance matrix and the multivariate normal distribution (known for students who have taken TMA4267 Linear statistical models).

What will you learn?

- Statistical learning and examples thereof
- Introduce relevant notation and terminology
- Prediction accuracy vs. model interpretability
- Bias-variance trade-off
- The basics of random vectors, covariance matrix and the multivariate normal distribution.

What is statistical learning?

Statistical learning is the process of learning from data.

By applying *statistical methods* on a *data set* (called the *training set*), we would like to *draw conclusions* about the relations between the variables (aka *inference*) or to *find a predictive function* (aka *prediction*) for new observations.

Also, we would like to find structures in the data that help us to learn something about the real world.

Statistical learning plays a key role in many areas of science, finance and industry.

Key concepts (stats) and notation

Todo (see M1L1notes.pdf)

Two variable types

Quantitative variables are variables from a continuous set, they have a numerical value.

- Examples: a person's weight, a company's income, the age of a building, the temperature outside, the level of precipitation etc.


Qualitative variables are variables from a discrete set, from a set of K different classes/labels/categories.

- Examples: type of fruit {apples, oranges, bananas, ...}, sex {male, female, other }, education level.
- Qualitative variables which have only two classes are called *binary* variables and are usually coded by 0 (no) and 1 (yes).

Examples of learning problems


- To predict the price of a stock 3 months from now, based on company performance measures and economic data. Here the response variable is quantitative (price).
- Spam detection for emails.
- To identify the risk factors for Prostate cancer.
- Estimating the risk that someone will suffer from a heart disease or heart attack, given knowledge about condition, behaviour, age, or demographic, diet and clinical measurements. Here the *outcome is binary* (yes, no) with both qualitative and quantitative input variables.
- Predict someone's body fat, given BMI, weight, age, etc.
- Digit recognition.

South African coronary heart disease data: 462 observations and 10 variables.



Handwritten digit recognition

To identify the numbers in a handwritten ZIP code, from a digitized image. This is a classification problem, where the response variable is categorical with classes $\{0, 1, 2, \dots, 9\}$ and the task is to correctly predict the class membership.



Examples of handwritten digits from U.S. postal envelopes.

Image taken from <https://web.stanford.edu/~hastie/ElemStatLearnII/>

Email classification (spam detection)

- Goal: to build a spam filter.
- This filter can be based on the frequencies of words and characters in emails.

The table below shows the average percentage of words or characters in an email message, based on 4601 emails of which 1813 were classified as a spam.

| | you | free | george | ! | \$ | edu |
|----------|------|------|--------|------|------|------|
| not spam | 1.27 | 0.07 | 1.27 | 0.11 | 0.01 | 0.29 |
| spam | 2.26 | 0.52 | 0.00 | 0.51 | 0.17 | 0.01 |


What makes a Nobel Prize winner?

Perseverance, luck, skilled mentors or simply chocolate consumption? An article published in the New England Journal of Medicine have concluded with the following:

Chocolate consumption enhances cognitive function, which is a sine qua non for winning the Nobel Prize, and it closely correlates with the number of Nobel laureates in each country. It remains to be determined whether the consumption of chocolate is the underlying mechanism for the observed association with improved cognitive function.

The figure shows the number of Nobel Laureates per 10 million population against countries' annual per capita chocolate consumption.

You can read the article [here](#) and a informal review of the article [here](#).



The Supervised Learning Problem

Starting point:

- Outcome measurement Y (also called dependent variable, response, target).
- Vector of p predictor measurements X (also called inputs, regressors, covariates, features, independent variables).
- In the **regression problem**, Y is quantitative (e.g price, blood pressure).
- In the **classification problem**, Y takes values in a finite, unordered set (survived/died, digit 0-9, cancer class of tissue sample).
- We have training data $(x_1, y_1), \dots, (x_N, y_N)$. These are observations (examples, instances) of these measurements.

Supervised learning and its objectives

Our data set (training set) consists of n measurement of the response variable Y and of p covariates x :

$$(y_1, x_{11}, x_{12}, \dots, x_{1p}), (y_2, x_{21}, \dots, x_{2p}), \dots, (y_n, x_{n1}, x_{n2}, \dots, x_{np}).$$

On the basis of the *training data* we would like to:

- **Accurately predict** unseen test cases.
- **Understand** which input affects the outcomes, and how.
- **Assess the quality** of your predictions and inference.

Supervised learning examples (we will study):

- Linear regression (M3), Logistic regression (M4), Generalized additive models (M7)
- Classification trees, bagging, boosting (M8), K-nearest neighbor classifier (M4)
- Support vector machines (M9)
- Neural networks (M11)

→ The majority of problems studied here fall in the supervised learning category

The Unsupervised Learning Problem

- There is **no outcome variable** y , just a set of predictors (features) x_i measured on a set of samples.
- Objective is more fuzzy – find (hidden) patterns or groupings in the data - in order to *gain insight and understanding*. There is no *correct* answer.
- Difficult to know how well you are doing.
- Different from supervised learning, but can be useful as a pre-processing step for supervised learning.

Examples in the course:

- Clustering (M10)
- Principal component analysis (M10)

Semi-supervised learning

(Not considered in this course)

- Our data set consists of a input data, and some of the data has labelled responses.
- This situation can for example occur if the measurement of input data is cheap, while the output data is expensive to collect.
- Classical solutions (likelihood-based) to this problem exists in statistics (missing at random observations).

Task

Find examples to explain the difference between supervised and unsupervised learning.

-
-
-

Overall philosophy

It is important to understand the ideas behind the various techniques, in order to know how and when to use them.

- One has to understand the simpler methods first, in order to grasp the more sophisticated ones.
- It is important to accurately assess the performance of a method, to know how well or how badly it is working

→ **Simpler methods often perform as well as fancier ones!**

- Statistical learning is a fundamental ingredient in the training of a modern data scientist.
- It is an exciting research area, having important applications in science, industry and finance.

Statistical Learning vs. Machine Learning

- Machine learning arose as a subfield of Artificial Intelligence.
- Statistical learning arose as a subfield of Statistics.
- There is much overlap — both fields focus on supervised and unsupervised problems:
 - Machine learning has a greater emphasis on large scale applications and prediction accuracy.
 - Statistical learning emphasizes models and their interpretability, and precision and uncertainty.
- The distinction has become more and more blurred, and there is a great deal of “cross-fertilization”.
- Machine learning has the upper hand in Marketing!

- There is a controversy and some scepticism against “too fancy” ML methods.
- Criticism: ML often re-invents existing methods and names them differently, but often without awareness of existing methods in statistics.
- Almost weekly new literature that delivers comparison. Often, the “simple” statistical methods “win”.

RESEARCH ARTICLE

Statistical and Machine Learning forecasting methods: Concerns and ways forward

Spyros Makridakis¹, Evangelos Spiliotis^{2*}, Vassilios Assimakopoulos²

1 Institute For the Future (IFF), University of Nicosia, Nicosia, Cyprus, 2 Forecasting and Strategy Unit, School of Electrical and Computer Engineering, National Technical University of Athens, Zografou, Greece

* spiliotis@fsu.gr

What is the aim in statistical learning?

Assume:

- we observe one *quantitative* response Y and
- p different predictors x_1, x_2, \dots, x_p .

We assume that there is a function f that relates the response and the predictor variables:




$$Y = f(x) + \varepsilon,$$

where ε is a random error term with mean 0 and independent of x .

The aim is to estimate f .



Example 1

Sales of a product, given advertising budgets in different media.



Example 2

Income for given levels of education.



There are two main reasons for estimating f :

- **Prediction**
- **Inference**

Reason 1: Prediction

Aim: predict a response Y given new observations x of the covariates as accurately as possible.

Notation:

$$\hat{Y} = \hat{f}(x).$$

- \hat{f} : estimated f
- \hat{Y} prediction for Y given x .
- We do not really care about the shape of f (“black box”).
→ no interpretation of regression parameters when the aim is purely prediction!

There are two quantities which influence the accuracy of \hat{Y} as a prediction of Y :

- The *reducible error* has to do with our estimate \hat{f} of f . This error can be reduced by using the most *appropriate* statistical learning technique.
- The *irreducible error* comes from the error term ε and cannot be reduced by improving f . This is related to the unobserved quantities influencing the response and possibly the randomness of the situation.

For a given \hat{f} and a set of predictors X which gives $\hat{Y} = \hat{f}(X)$, we have

$$\mathbb{E}[(Y - \hat{Y})^2] = \underbrace{\mathbb{E}[(f(X) - \hat{f}(X))^2]}_{\text{reducible}} + \underbrace{\text{Var}(\epsilon)}_{\text{irreducible}}$$

Q: If there were a *deterministic* relationship between the response and a set of predictors, would there then be both reducible and irreducible error?

Reason 2: Inference

Aim: *understand* how the response variable is affected by the various predictors (covariates).

The *exact form* of \hat{f} is of *main interest*.

- Which predictors are associated with the response?
- What is the relationship between the response and each predictor?
- Can the relationship be linear, or is a more complex model needed?

Estimating f

Overall idea:

- Using available *training data* $(x_1, y_1), \dots, (x_n, y_n)$ to estimate \hat{f} , such that $Y \approx \hat{f}(X)$ for any (X, Y) (also those that have not yet been observed).

Two main approaches:

- Parametric methods
- Non-parametric methods

Parametric methods

Assumption about the form or shape of the function f .

The multiple linear model (M3) is an example of a parametric method:

$$f(x) = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p + \varepsilon,$$

with $\varepsilon \sim N(0, \sigma^2)$.

The task simplifies to finding estimates of the $p + 1$ coefficients $\beta_0, \beta_1, \dots, \beta_p$. To do this we use the training data to fit the model, such that

$$Y \approx \hat{\beta}_0 + \hat{\beta}_1 x_1 + \dots + \hat{\beta}_p x_p.$$

Fitting a parametric models is thus done in two steps:

1. Select a form for the function f .
2. Estimate the unknown parameters in f using the training set.

Non-parametric methods

- Non-parametric methods seek an estimate of f that gets close to the data points, but without making explicit assumptions about the form of the function f .
- The K -nearest neighbour algorithm is an example of a non-parametric model. Used in classification, this algorithm predicts a class membership for a new observation by making a majority vote based on its K nearest neighbours. We will discuss the K -nearest neighbour algorithm later in this module.

Q: What are advantages and disadvantages of parametric and non-parametric methods?

Hints: interpretability, amount of data needed, complexity, assumptions made, prediction accuracy, computational complexity, over/under-fit.

A: Parametric methods

| Advantages | Disadvantages |
|--------------------------------------|---|
| Simple to use and easy to understand | The function f is constrained to the specified form. |
| Requires little training data | The assumed function form of f will in general not match the true function, potentially giving a poor estimate. |
| Computationally cheap | Limited flexibility |

A: Non-parametric methods

| Advantages | Disadvantages |
|--|--|
| Flexible: a large number of functional forms can be fitted | Can overfit the data |
| No strong assumptions about the underlying function are made | Computationally more expensive as more parameters need to be estimated |
| Can often give good predictions | Much data is required to estimate (the complex) f . |

Prediction accuracy vs. interpretability

(we are warming up to the bias–variance trade–off)

Inflexible methods:



- Linear regression (M3)
- Linear discriminant analysis (M4)
- Subset selection and lasso (M6)

Flexible methods:

- KNN classification (M2, M4), KNN regression, Smoothing splines (M7)
- Bagging and boosting (M8), support vector machines (M9)
- Neural networks (M11)

Why would I ever prefer an inflexible method?

Example: Prediction of income from “Years of Education” and “Seniority”



The choice of a flexible or inflexible method depends on the goal in mind.

- For **inference** an inflexible model is easier to interpret.
- For **prediction** a flexible model is more powerful.


Overfitting occurs when the estimated function f is too closely fit to the observed data points.

Underfitting occurs when the estimated function f is too rigid to capture the underlying structure of the data.

We illustrate this by a toy example using polynomial regression.

Polynomial regression example (simulation)

Consider a covariate x observed between -2 to 4 and $n = 61$ observations.




We impose a relationship between response Y and covariate x :

$$Y = x^2 + \varepsilon$$

with error (noise) term $\varepsilon \sim N(0, \sigma^2)$ with $\sigma = 2$. It is a substitute for all the unobserved variables that are not in our equation, but that might influence Y .

We call $Y = x^2$ the *truth*.





Truth



Next, we want to fit a function to the observations *without* knowing the true relationship, and we have tried different parametric polynomial functions.

- **poly1:** Simple linear model of the form $\beta_0 + \beta_1 x$ fitted to the observations. *Underfits* the data.
- **poly2:** Quadratic polynomial fit to the data, of the form $\beta_0 + \beta_1 x + \beta_2 x^2$. This fits well.
- **poly10:** Polynomial of degree 10 fit of the form $\beta_0 + \beta_1 x + \beta_2 x^2 + \cdots + \beta_{10} x^{10}$ *Overfits* the data.
- **poly20:** Polynomial of degree 10 fit of the form $\beta_0 + \beta_1 x + \beta_2 x^2 + \cdots + \beta_{20} x^{20}$ *Overfits* the data.

We will discuss polynomial regression in M7.



The degree of the polynomial is a *flexibility parameter*.

We can now ask:

- Which of these models performs “best”?
- Is there *one* method that dominates all others?

Assessing model accuracy

No method dominates all others over all possible data sets.

- That is why we need to learn about many different methods.
- For a given data set we need to know how to decide which method produces the *best* results.
- We need to understand what *best* means.
- How close is the predicted response to the true response value?

Measuring the Quality of Fit

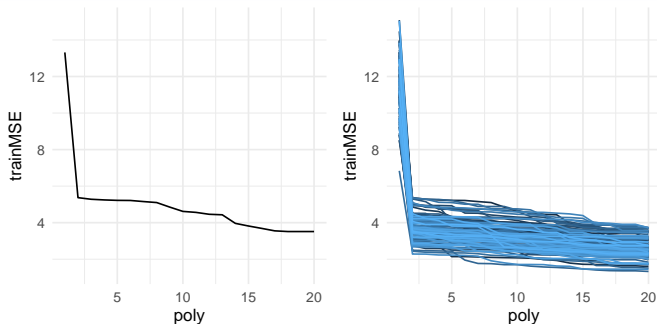
- A popular measure for the quality of fit: *Training MSE* (mean squared error).
- It is the mean of squared differences between prediction and truth for the training data (the same values that were used to estimate f):

$$\text{MSE}_{\text{train}} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2$$

- But: we are *not* interested in how the method works on the training data. We want to know how good the method is when we use it on *previously unseen test data* (e.g., future data).

Examples:

- We don't want to predict last weeks stock price, we want to predict the stock price next week.
- We don't want to predict if a patient in the training data has diabetes (because we already know this), we want to predict if a new patient has diabetes.



Q: Based on the training MSE - which model fits the data the best?

Polynomial example: fitted order 1-20 polynomial when the truth is order 2. Left: one repetition, right: 100 repetitions of the training set.

Test MSE

- Simple solution: estimate \hat{f} using the training data (maybe by minimizing the training MSE), but choose the *best* model using a separate *test set*.
- *Test MSE* for a set of n_0 test observations (x_{0j}, y_{0j}) :

$$\text{MSE}_{\text{test}} = \frac{1}{n_0} \sum_{j=1}^{n_0} (y_{0j} - \hat{f}(x_{0j}))^2$$

- Alternative notation:

$$\text{Ave}(y_0 - \hat{f}(x_0))^2$$



(taking the average over all available test observations).

Q: What if we do not have access to test data?

A: In Module 5 we will look into using *cross validation* to mimic the use of a test set.





Q: But, can we instead just use the training data MSE to choose a model? A low training error should also give a low test error?

A: Sadly no, if we use a flexible model we will look at several cases where a low training error is a sign of overfitting, and will give a high test error. So, the training error is not a good estimator for the test error because it does not properly account for model complexity.



Polynomial example: fitted order 1-20 when the truth is order 2.
Left: one repetition, right: 100 repetitions for the testMSE.

Q: Based on the test MSE - which model fits the data the best?



A: If choosing flexibility based on training MSE=poly20 wins, if choose flexibility based on test MSE=poly 2 wins.

Test error vs. training error

Important observations:

- The test error seems to have a minimum (U-shape) in between the extremes.
- The training error keeps going down.

Why?

The Bias-Variance trade-off

- Assume we have fitted a *regression* curve $Y = f(x) + \varepsilon$ to our training data, which consist of independent observation pairs $\{x_i, y_i\}$ for $i = 1, \dots, n$. (Yes, only one covariate x .)
- Assume that ε is an unobserved random variable that adds noise to the relationship between the response variable and the covariates (random error), with mean zero and constant variance σ^2 for all values of x .
- ε is a substitute for all the unobserved variables that influence Y .
- The training data was used to estimate \hat{f} .

- We want to use \hat{f} to obtain the predicted response value $\hat{f}(x_0)$ for an unseed test observation (x_0, y_0) .
- The *expected test mean squared error (MSE)* at x_0 is defined as:

$$E[y_0 - \hat{f}(x_0)]^2$$

- Compare this to the test MSE for the polynomial example (MSE_{test}): The average is simply replaced by the *theoretical version* (expected value).

Using that $y_0 = f(x_0) + \varepsilon$, this expected test MSE can be decomposed into three terms

$$\begin{aligned} & \mathbb{E}[f(x_0) + \varepsilon - \hat{f}(x_0)]^2 \\ &= \mathbb{E}[f(x_0)^2] + \mathbb{E}[\varepsilon^2] + \underbrace{\mathbb{E}[\hat{f}(x_0)^2]}_{=f(x_0)^2} - 2\mathbb{E}[f(x_0)\hat{f}(x_0)] \\ &= \text{Var}(f(x_0)) + \mathbb{E}[f(x_0)^2] + \text{Var}(\varepsilon) + \text{Var}(\hat{f}(x_0)) + \mathbb{E}[\hat{f}(x_0)]^2 \\ &\quad - 2f(x_0)\mathbb{E}[\hat{f}(x_0)] \\ &= \underbrace{\text{Var}(\varepsilon)}_{\text{Irreducible error}} + \underbrace{\text{Var}(\hat{f}(x_0))}_{\text{Variance of prediction}} + \underbrace{\left(f(x_0) - \mathbb{E}[\hat{f}(x_0)]\right)^2}_{\text{Squared bias}} \end{aligned}$$

Q: what assumptions have we made in the derivation above?

A: classnotes.

$$E[(y_0 - \hat{f}(x_0))^2] = \dots = \text{Var}(\varepsilon) + \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2$$

- First term: irreducible error. This term cannot be reduced regardless how well our statistical model fits the data.
- Second term: variance of the prediction at x_0 or the expected deviation around the mean at x_0 . If the variance is high, there is large uncertainty associated with the prediction.
- Third term: squared bias. The bias gives an estimate of how much the prediction differs from the true mean. If the bias is low the model gives a prediction which is close to the true value.

Note:

$$E[(y_0 - \hat{f}(x_0))^2]$$





is the **expected test MSE**. We can think of this as the average test MSE we would obtain if we repeatedly estimated f using many training sets (as we did in our example), and then tested this estimate at x_0 .

However, if we also assume that X is a random variable, this is really $E[(Y - \hat{f}(x_0))^2 \mid X = x_0]$

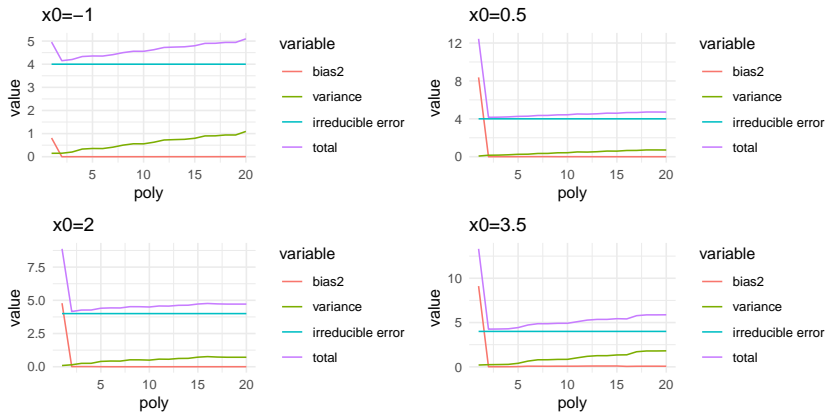
The **overall expected test MSE** can we then compute by averaging the expected test MSE over all possible values of x_0 (averaging with respect to frequency in test set), or mathematically by the law of total expectation $E\{E[(Y - \hat{f}(X))^2 \mid X]\}$ (also sometimes referred to as the law of double expectations).

Polynomial example (cont.)

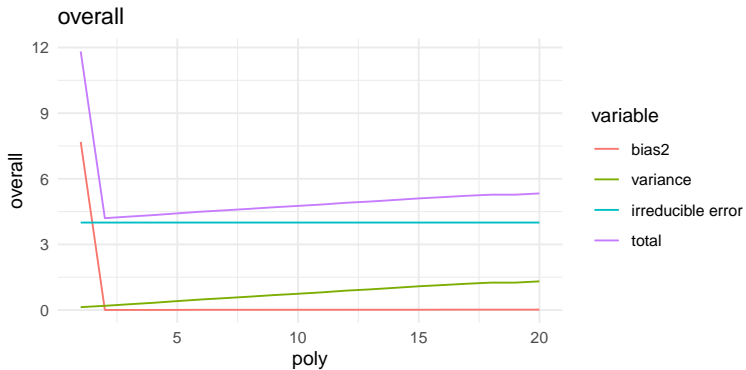
$Y = x^2$ is still the *truth*.



For 4 different polynomial models (poly1,2,10 and 20), the squared bias, variance, irreducible error and the total sum. Plots based on 100 simulations for the polynomial example.



At 4 different values for x_0 , the squared bias, variance, irreducible error and the total sum. Plots based on 100 simulations for the polynomial example.






Overall version (averaging over 61 gridpoints of x).

Choosing the best model

When fitting a statistical model the aim is often to obtain **the most predictive model**.

- **Training set:** The observations used to fit the statistical model → Training error
- **Test sample:** new observations which were not used when fitting the model → Test error
- Training error decreases for more complex models, but the test error has an **optimum**.
- This trade-off in selecting a model with the right amount of complexity/flexibility is the **Bias-Variance trade-off**.



- Inflexible models may lead to a poor fit (high bias)
- Flexible (complex) models may provide more unbiased fits but may overfit the data (high variance)
- The aim is to find the optimum.

References

James, Gareth, Daniela Witten, Trevor Hastie, and Robert Tibshirani. 2013. *An Introduction to Statistical Learning*. Vol. 112. Springer.