

TMA4268 Statistical Learning

Chapter 10: Unsupervised Learning - Lab 1

Thiago G. Martins, Department of Mathematical Sciences, NTNU

Spring 2020

Lab 1: Principal Components Analysis

We are going to use the `USArrests` dataset. It contains crime statistics per 100000 residents in 50 states of the USA. The crime types included are Assault, Murder and Rape. In addition, it contains information about the percent of the population in each state living in urban areas.

Data Exploration

Data rows

The rows of the data contains the 50 states, in alphabetical order.

```
states=row.names(USArrests) # part of the base R package
states
```

```
## [1] "Alabama"      "Alaska"       "Arizona"      "Arkansas"
## [5] "California"   "Colorado"     "Connecticut"  "Delaware"
## [9] "Florida"     "Georgia"      "Hawaii"       "Idaho"
## [13] "Illinois"     "Indiana"      "Iowa"         "Kansas"
## [17] "Kentucky"    "Louisiana"    "Maine"        "Maryland"
## [21] "Massachusetts" "Michigan"     "Minnesota"    "Mississippi"
## [25] "Missouri"    "Montana"      "Nebraska"     "Nevada"
## [29] "New Hampshire" "New Jersey"   "New Mexico"   "New York"
## [33] "North Carolina" "North Dakota" "Ohio"         "Oklahoma"
## [37] "Oregon"      "Pennsylvania" "Rhode Island" "South Carolina"
## [41] "South Dakota" "Tennessee"    "Texas"        "Utah"
## [45] "Vermont"     "Virginia"     "Washington"   "West Virginia"
## [49] "Wisconsin"   "Wyoming"
```

Data columns

The columns of the data contain the four variables

```
names(USArrests)
```

```
## [1] "Murder" "Assault" "UrbanPop" "Rape"
```

Mean and variance

The variables have vastly different means

```
apply(USArrests, 2, mean)
```

```
## Murder Assault UrbanPop Rape
## 7.788 170.760 65.540 21.232
```

The same is true for the variances

```
apply(USArrests, 2, var)
```

```
## Murder Assault UrbanPop Rape
## 18.97047 6945.16571 209.51878 87.72916
```

If we failed to scale the variables before performing PCA, then most of the principal components that we observed would be driven by the Assault variable, since it has by far the largest mean and variance.

PCA

We now go on to apply PCA on the standardized variables (mean 0 and std. dev. 1). By default, the `prcomp()` function centers the variables to have mean zero. By using the option `scale=TRUE`, we scale the variables to have standard deviation one.

```
pr.out=prcomp(USArrests, scale=TRUE)
```

PCA output

The output from `prcomp()` contains a number of useful quantities.

```
names(pr.out)
```

```
## [1] "sdev" "rotation" "center" "scale" "x"
```

center and scale

The `center` and `scale` components correspond to the means and standard deviations of the variables that were used for scaling prior to implementing PCA.

```
pr.out$center
```

```
## Murder Assault UrbanPop Rape
## 7.788 170.760 65.540 21.232
```

```
pr.out$scale
```

```
## Murder Assault UrbanPop Rape
## 4.355510 83.337661 14.474763 9.366385
```

rotation

The rotation matrix provides the principal component loadings; each column of `pr.out$rotation` contains the corresponding principal component loading vector.

```
pr.out$rotation
```

```
## PC1 PC2 PC3 PC4
## Murder -0.5358995 0.4181809 -0.3412327 0.64922780
## Assault -0.5831836 0.1879856 -0.2681484 -0.74340748
## UrbanPop -0.2781909 -0.8728062 -0.3780158 0.13387773
## Rape -0.5434321 -0.1673186 0.8177779 0.08902432
```

Score vectors

The 50×4 matrix `x` has as its columns the principal component score vectors. That is, the k th column is the k -th principal component score vector.


```
dim(pr.out$x)
```

```
## [1] 50 4
```

Biplot

We can plot the first two principal components as follows:


```
biplot(pr.out, scale=0)
```



The `scale=0` argument to `biplot()` ensures that the arrows are scaled to represent the loadings; other values for scale give slightly different biplots with different interpretations.

Remember that the PCs are unique up to a sign change, so the code below should give equivalent results

```
pr.out$rotation=-pr.out$rotation  
pr.out$x=-pr.out$x  
biplot(pr.out, scale=0)
```



Proportion of variance explained (PVE)

The `prcomp()` function also outputs the standard deviation of each principal component.

```
pr.out$sdev
```

```
## [1] 1.5748783 0.9948694 0.5971291 0.4164494
```

The variance explained by each principal component is obtained by squaring these:

```
pr.var=pr.out$sdev^2
pr.var
```

```
## [1] 2.4802416 0.9897652 0.3565632 0.1734301
```


We can then compute the proportion of variance explained by each principal component

```
pve=pr.var/sum(pr.var)
pve
```

```
## [1] 0.62006039 0.24744129 0.08914080 0.04335752
```

We can plot the PVE explained by each component, as well as the cumulative PVE, as follows:

```
plot(pve,  
     xlab="Principal Component",  
     ylab="Proportion of Variance Explained",  
     ylim=c(0,1),  
     type='b')
```



```
plot(cumsum(pve),  
     xlab="Principal Component",  
     ylab="Cumulative Proportion of Variance Explained",  
     ylim=c(0,1),  
     type='b')
```

